

# Container Security

Presented by:

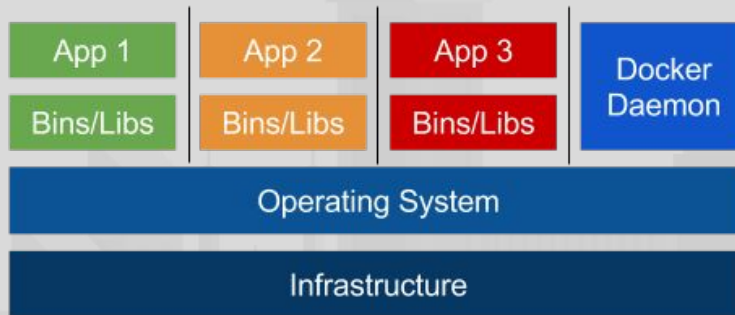
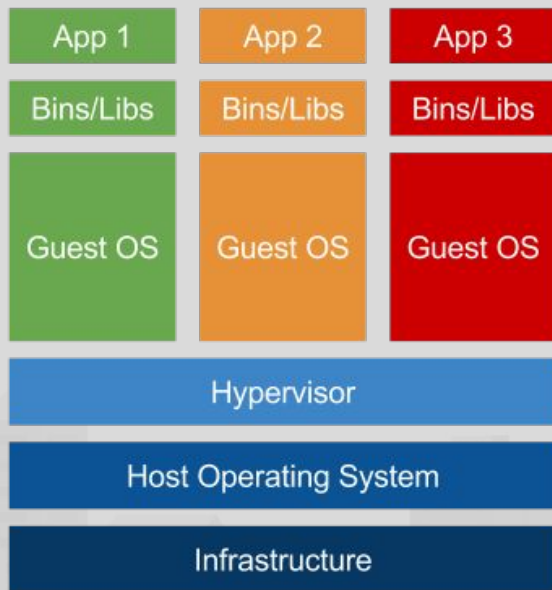
Kasra Motamedi (Security Engineer)  
Amir Moghaddam (Security Engineer)

Spring 1401

# What Are Containers?

- Resolving Dependency Nightmare
- Build once, run anywhere
- Simple, fast, and transparent way to set up runtime environments
- Easy for developers and operations to understand and use
- Smaller runtime footprint than VMs, fast deploy and startup
- Ideal for microservices

# Containers vs. VMs



# Container Advantages (Security Perspective)

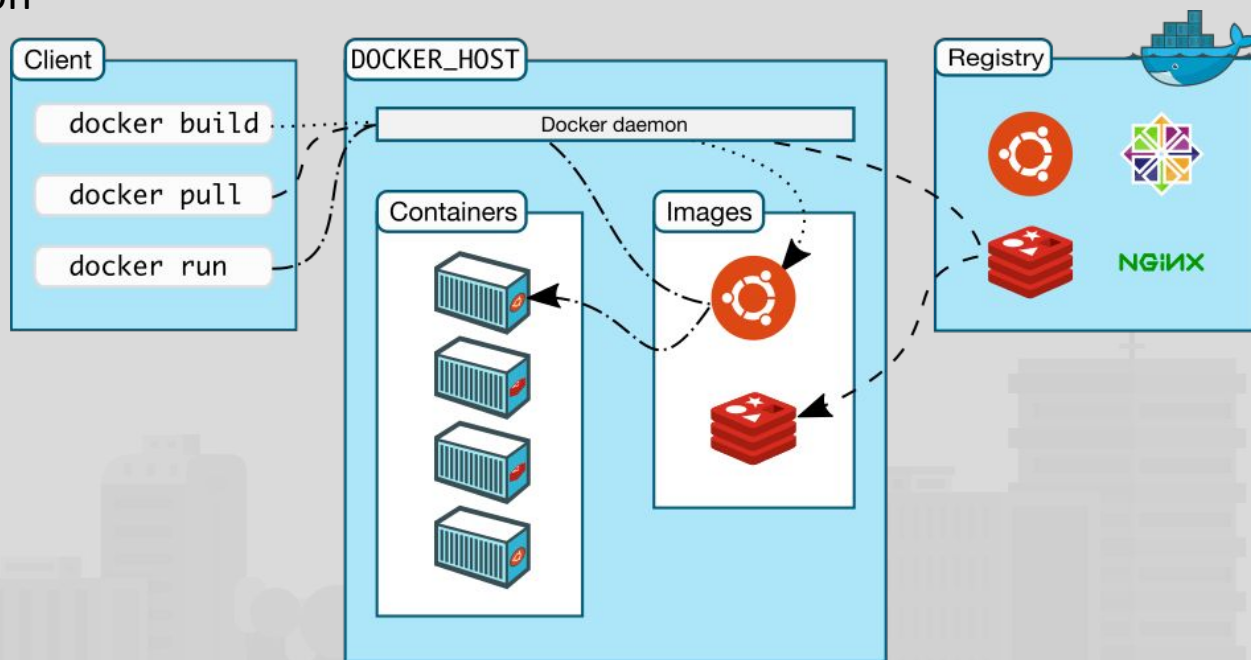
- Runtime Separation of applications **on a host**
- Only deploy what you need: limited dependencies and smaller attack surface
- Easier to set up and tear down instances makes it easier to patch
- Increases consistency
  - making hardening, auditing, and testing easier
- Docker runtime has “sane OS defaults”: limited capabilities, etc.

# Container Security Issues


- Docker Daemon presents its own attack surface
- Managing secrets with Docker
- Lightweight isolation
- Process running as root in container has root in underlying host
- Untrusted content
  - compromised and vulnerable images
- Container sprawl, especially at scale, and ephemeral runtimes are difficult to track and manage using classic network defenses (WAFs, IDS/IPS...)

# Docker Components


- Host/Docker Daemon
- Registry
- Dockerfile
- Image
- Container



# Image Attack Vectors



# Container Attack Vectors





# Host/Docker Daemon Security

- Monitoring host vulnerabilities
- Checking members of docker group
- Set correct permissions for docker files
- Track and patch vulnerabilities in Docker engine
- Enabling docker content trust (DCT)
  - Export `DOCKER_CONTENT_TRUST=1`
  - Image tags are digitally signed before pushing to a registry
  - Enforce using digitally signed images (push and pull)
- Enable auditing docker files
- Tool → Docker Bench Security [[demo-link](#)]
  - Scanning docker daemon + images + containers for best practices

# Host/Docker Daemon Security (Authentication/Authorization)

- Do not expose the daemon API/socket unless you enable mutual TLS and user authentication
- If you have socket access to Docker Daemon, you have host root [[demo-link](#)]
- Authorization plugin framework implemented in Docker 1.10 so that you can write your own checks/rules
- Tool: Twistlock AuthZ Broker

# Registry Security

- Using private docker registry
  - Preventing access to app codes
  - Part of CI/CD process
  - Faster pull/push
- Changing default registry port
- Enable TLS support and authentication

# Image Security (1)

- Images from public registries can contain vulnerabilities or malware
- Use official images from public repositories
- Keeping images up to date
  - Protecting against vulnerabilities
  - Microsoft fixed vulnerabilities 2 months before WannaCry attack
- Scanning images regularly
  - Detecting vulnerabilities

# Image Security (2)

- Signing images
  - Integrity check
- Don't use setuid
  - Set the least capabilities
- Tool → Trivy [[demo-link](#)]
  - Auditing images for best practices
- Tool → Dive [[demo-link](#)]
  - Analyzing each image layer contents
  - Searching files in each layer for forensic purposes

# Dockerfile Best Practices (1)

- Use base image from trusted sources
- Use scratch images
  - Reduce the attack surface
- Copy only what is needed
- Use non root user account
- Manage secrets with Swarm or Kubernetes if you are running in clusters
- Docker build secrets (--secrets) switch released in 18.09
- Best practice is to use a general-purpose secret keeper like Vault
- Don't use secrets in Dockerfile [[demo-link](#)]

# Dockerfile Best Practices (2)

- Set least privileges for non root users
  - If it doesn't need, only set read and (or) execute for non root users
- Don't run unnecessary services inside the Dockerfile
  - Reduce the attack surface
- Use multi stage builds
  - Reduce the attack surface
- Tool → Trivy [[demo-link](#)]
  - Auditing Dockerfiles for best practices

# Container Security (Isolation+Limitation)

- Isolation through (primarily) kernel namespaces
- Vulnerabilities in host OS kernel or Docker Engine can break isolation
- Config mistakes can allow processes in containers to interact with other containers and with the host
- Use selective sharing of namespaces between containers
- Don't use host Namespaces (e.x., --uid host)
- Use Cgroups
  - Resource limitation (cpu and memory)
  - Prevent DOS attacks

```
root@ubuntu:~# lsns -p 1
      NS TYPE      NPROCS  PID  USER  COMMAND
4026531835 cgroup      350    1  root  /sbin/init splash
4026531836 pid        339    1  root  /sbin/init splash
4026531837 user       333    1  root  /sbin/init splash
4026531838 uts        347    1  root  /sbin/init splash
4026531839 ipc        345    1  root  /sbin/init splash
4026531840 mnt         327    1  root  /sbin/init splash
4026532008 net        334    1  root  /sbin/init splash
root@ubuntu:~#
```



# Container Security (User Namespace)

- Prevent using root user inside the containers
- Users are not fully namespaced in Docker
- User namespace remapping at runtime added in Docker 1.10
  - `--userns-remap=[USERNAME]`
  - NOT enabled by default

# Container Security (Capabilities)

- Default limited capabilities in Docker → reducing the container attack service
- Dropping additional capabilities
  - Note: `--cap-add` can be used to increase privileges if required
  - Best practice is to drop all capabilities not required by the application
  - [\[demo-link\]](#)
- Watch out: `--privileged` argument allows a Docker container to access all devices on the host and overrides AppArmor/SELinux restrictions! This should never be allowed
  - [\[demo-link\]](#)
  - [\[demo-link\]](#)

# Container Security (Capabilities)

## ➤ Capabilities

- CAP\_SYS\_PTRACE
  - process injection → reverse shell
- CAP\_SYS\_MODULE
  - kernel module injection → reverse shell
- CAP\_DAC\_READ\_SEARCH
  - read host files
- CAP\_DAC\_OVERRIDE
  - read and write host files
- CAP\_SYS\_TIME
  - set system clock

# Container Security (Docker Hosts)

- Some host OS builds are designed specifically to support Docker or the containers, especially in the cloud
  - CoreOS
  - RancherOS
  - Snappy Ubuntu Core
  - VMware Photon OS

# Container Security (Security Profiles/Policies)

## ➤ Seccomp

- Secure Computing Mode (seccomp)
- syscall firewall between user-level processes and the Linux Kernel
- Support and a default profile is implemented as of Docker Engine 1.10
- Whitelist or blacklist filters define that system calls and arguments are allowed
- Default filter blocks over 50 syscalls when a container runs as non-privileged user

## ➤ AppArmor

- Used for debian based distros
- Blocking capabilities for programs

## ➤ SELinux

- Used for redhat based distros
- Defining access controls for apps, processes and files

# Container Security (Other)

- Least privileges
  - Set the least capabilities if needed
- Do NOT pass secrets in runtime environment variables
- Mounting the container filesystem as read-only
- Not mounting sensitive host directories inside the containers
- Restricting privileges
  - `--security-opt=no-new-privileges` (disable setuid)
- Collect log files
  - `/var/lib/docker/containers/*//*.log`
  - `/etc/falco/events.txt`
- Tool → Falco [[demo-link](#)]
  - Detecting anomalous activity

# Summary

- Containers **can** improve security over bare metal/VM runtimes
- Don't depend on Docker defaults
  - Docker provides some built-in architectural sandboxing and other runtime protection, but is NOT secure by default
- Make appropriate trade-offs between developer flexibility and security
- Treat images as sensitive data
  - control and restrict access to images, review and scan image contents
- Major investments are being made in container security
  - expect improvements from Docker, the community, and vendors

# Reference

- <https://www.sans.org/cyber-security-courses/cloud-security-devsecops-automation/>
- <https://info.aquasec.com/container-security-book>
- <https://bootcamps.pentesteracademy.com/course/container-security-on-demand>
- <https://www.udemy.com/course/docker-security>
- [https://www.youtube.com/watch?v=b\\_euX\\_M82uI](https://www.youtube.com/watch?v=b_euX_M82uI)